

Monte Carlo Simulation

[Monte Carlo methods](#), or Monte Carlo experiments, are a broad class of computational algorithms that rely on repeated random sampling to obtain numerical results. The underlying idea is the weak law of large numbers.

Let X_1, X_2, \dots be i.i.d. with $E|X_i| < \infty$. Let $S_n = X_1 + \dots + X_n$ and let $\mu = EX_1$. Then $S_n/n \rightarrow \mu$ in probability.

Loosely speaking, **the expectation** can be well approximated by **the sample analogue** in probability as sample size n goes to infinity.

We will apply the monte carlo simulation to compute the expectations of interest.

- $Eh(X) = \int h(x)f(x)dx$ is well approximated by $\frac{1}{n} \sum_{i=1}^n h(X_i)$ as n goes to infinity, where X_1, \dots, X_n are i.i.d. with common pdf $f(x)$.
- Similarly, $Pr(X \leq x)$ is well approximated by $\frac{1}{n} \sum_{i=1}^n 1\{X_i \leq x\}$ as n goes to infinity.
 - Probabilities can be expressed as **expectations of indicator functions**, i.e.,

$$Pr(X \leq x) = \int_{-\infty}^x f(t)dt = \int_{-\infty}^{\infty} 1_{(-\infty, x]}(t)f(t)dt = E1_{(-\infty, x]}.$$

Example: $Pr(4 < X \leq 5)$

Consider the random variable $X \sim N(1, 4)$. We are interested in $Pr(4 \leq X < 5)$.

The `norm.cdf()` function can be used to compute the target probability directly.

```
In [1]: from scipy.stats import norm
import numpy as np
import matplotlib.pyplot as plt
mu = 1
sigma = 2
print("The probability that 4 < X <= 5 is %.3f" % (norm.cdf(5, mu, sigma) - norm
```

The probability that 4 < X <= 5 is 0.044

We can also get the approximate probability through Monte Carlo simulation.

First, generate 5000 random draws from the target distribution, i.e.,

$X_1, X_2, \dots, X_{5000} \sim N(1, 4)$.

```
In [2]: np.random.seed(123) # for reproducible randomness
nsim = 5000
```


```
X = norm.rvs(loc = mu, scale = sigma, size = nsim)
print("The first 5 random draws are ", X[:5])
```

The first 5 random draws are [-1.17126121 2.99469089 1.565957 -2.01258943 -0.1572005]

```
In [3]: np.random.seed(123) # given the same seed, we have the same random draws
        nsim = 5000
        norm.rvs(loc = mu, scale = sigma, size = nsim)[:5]
```

Out[3]: array([-1.17126121, 2.99469089, 1.565957 , -2.01258943, -0.1572005])

```
In [4]: plt.figure(figsize = (10, 5))
        plt.scatter(np.linspace(1, nsim, nsim), X, label = 'Random draws')
        plt.hlines(y = 5, xmin = 1, xmax = nsim, color = 'red', label = 'y = 5')
        plt.hlines(y = 4, xmin = 1, xmax = nsim, color = 'red', label = 'y = 4')
        plt.xlabel('Simulations')
        plt.legend()
        plt.show()
```



Now we compute the fraction of realizations (X_1, \dots, X_{5000}) falling between 4 and 5. This fraction is "approximately" the probability of interest when the number of runs is very large. (Weak law of large numbers)

```
In [5]: Y = (X > 4) * (X <= 5) * 1 # 5000 indicator variables
        print("The first five results: ", Y[:5])
        print("The fraction of X's satisfying the condition is %.3f" % (np.mean(Y) ))
```

The first five results: [0 0 0 0 0]

The fraction of X's satisfying the condition is 0.045

Note that 0.045 is quite close to the true probability 0.044 as computed above through `norm.cdf()` .

Example: $P(T \leq 0.6)$, where $T \sim t(5)$

Similarly, we use Monte Carlo simulation to conveniently approximate $P(T \leq 0.6)$, where $T \sim t(5)$.

In [6]:

```
from scipy.stats import t
x0 = 0.6
print('The probability that T <= {} is {:.3f}'.format(x0, t.cdf(x0, df = 5)))

np.random.seed(123)
nsim = 5000
X = t.rvs(df = 5, size = nsim)
print("The fraction of X's at most {} is {:.3f}".format(x0, np.mean(X <= x0)))
```

The probability that T <= 0.6 is 0.713
The fraction of X's at most 0.6 is 0.707

Example: $Pr(-1 \leq X_1 < 1, -2 \leq X_2 < 3)$

Suppose that

$$\begin{pmatrix} X_1 \\ X_2 \end{pmatrix} \sim N\left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{pmatrix} 1 & 0.5 \\ 0.5 & 4 \end{pmatrix}\right).$$

We can apply the monte carlo simulation to computing joint probabilities.

In [7]:

```
from scipy.stats import multivariate_normal
import pandas as pd
mu = np.array([0, 0])
cov = np.array([[1, 0.5], [0.5, 4]])
diff = multivariate_normal.cdf(np.array([1,3]), mean = mu, cov = cov) \
      - multivariate_normal.cdf(np.array([-1,3]), mean = mu, cov = cov) \
      - multivariate_normal.cdf(np.array([1,-2]), mean = mu, cov = cov) \
      + multivariate_normal.cdf(np.array([-1,-2]), mean = mu, cov = cov)
print("The target probability:", diff)

np.random.seed(123)
nsim = 50000
X = multivariate_normal.rvs(mean=mu, cov=cov, size=nsim)
result = (X[:, 0] < 1) * (X[:, 0] >= -1) * (X[:, 1] < 3) * (X[:, 1] >= -2) * 1
print("The fraction of random draws satisfying the condition is {:.3f}" % (np.sum(
```

The target probability: 0.5354662279197304
The fraction of random draws satisfying the condition is 0.535

In [8]:

```
result[:5]
```

Out[8]: array([0, 0, 0, 0, 1])

Example: $Pr(-3 \leq 2X_1 + 3X_2 < 3)$

Next we consider a new random variable $Y \equiv 2X_1 + 3X_2$. What is the probability $Pr(-3 \leq Y < 3)$?

Note that Y is normally distributed with mean 0 and variance

$2^2 * Var(X_1) + 3^2 * Var(X_2) + 2 * 2 * 3 * Cov(X_1, X_2) = 46$. So the true probability is

```
In [9]: print("The true probability is %.3f" % (norm.cdf(3, 0, np.sqrt(46)) - norm.cdf(-
```

The true probability is 0.342

We can also approximate the probability through Monte Carlo simulation.

```
In [10]: newvar = 2*X[:, 0] + 3*X[:,1]
fraction = np.mean((newvar >= -3) * (newvar < 3))
print("The fraction of random draws satisfying the condition is %.3f" % fraction
```

The fraction of random draws satisfying the condition is 0.342

```
In [11]: newvar
```

```
Out[11]: array([-5.76922309, -0.28063664, -1.47137681, ...,  0.98165658,
        7.8461253 , 13.25262412])
```

As we can see, the result from the monte carlo simulation is "close to" the true probability. In addition, **you do not need to derive the marginal distribution of the new random variable Y in the MC simulation.**

Example: $E[2X_1 + 3X_2]$ and $Var(2X_1 + 3X_2)$

As derived above, the mean and variance of $2X_1 + 3X_2$ are 0 and 46, respectively. This can be approximated through sample means and sample variance.

```
In [12]: print('The sample mean is %.3f' % np.mean(newvar))
print("The sample variance is %.3f" % np.var(newvar, ddof = 1))
```

The sample mean is 0.013

The sample variance is 46.313

Example: Approximating π

Suppose that $X_1, X_2 \stackrel{i.i.d.}{\sim} Uniform([-1, 1])$. We can draw a lot of samples from this joint distribution. As will be seen shortly, the area of the circle of radius 1 is π while the area of the square $[-1, 1] \times [-1, 1]$ is 4. Roughly, the fraction of the samples falling into the circle is $\pi/4$. Thus, we can infer the value of π through 4 times the fraction.

```
In [13]: np.pi
```

```
Out[13]: 3.141592653589793
```


In [14]:

```
from scipy.stats import uniform

nsim = 1000
X = uniform.rvs(loc = -1, scale = 2, size = (nsim, 2)) # the uniform distribution

fig = plt.figure(figsize = (6,6))
ax = fig.add_subplot(111)
plt.scatter(X[:, 0], X[:, 1])
circle = plt.Circle((0, 0), 1, color='r', fill=False)
ax.add_patch(circle)
plt.xlabel('x')
plt.ylabel('y')
plt.title('Number of simulations = {}'.format(nsim))
plt.show()

fraction = np.sum(X[:,0] ** 2 + X[:, 1] ** 2 <= 1) / nsim
print("The fraction of samples falling in the circle is %.3f" % fraction)
print("The estimated pi is %.3f" % (4*fraction))
```



The fraction of samples falling in the circle is 0.787
The estimated pi is 3.148

We can define a function of `nsim` and see the approximation errors in general decline with `nsim`.

In [15]:

```
def simulate_pi(nsim):
    X = uniform.rvs(loc = -1, scale = 2, size = (nsim, 2)) # the uniform distribution


    fig = plt.figure(figsize = (6,6))
    ax = fig.add_subplot(111)
    plt.scatter(X[:, 0], X[:, 1])
    circle = plt.Circle((0, 0), 1, color='r', fill=False)
    ax.add_patch(circle)
    plt.title('Number of simulations = {}'.format(nsim))
```

```
plt.xlabel('x')
plt.ylabel('y')
plt.show()

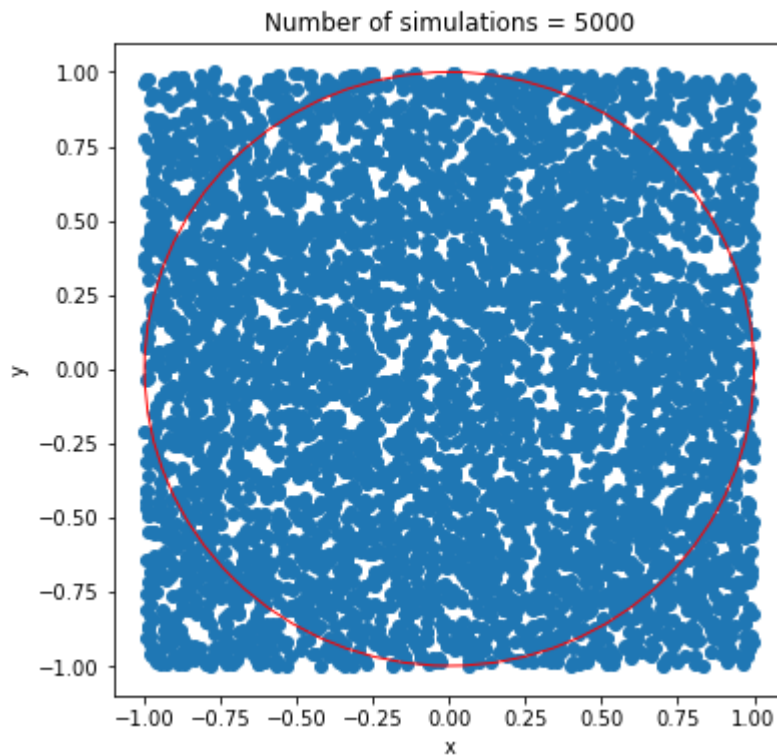
fraction = np.sum(X[:,0] ** 2 + X[:, 1] ** 2 <= 1) / nsim
print("The fraction of samples falling in the circle is %.3f" % fraction)
print("The estimated pi is %.3f" % (4*fraction))

np.random.seed(23)
simulate_pi(2000)

np.random.seed(123)
simulate_pi(5000)
```



The fraction of samples falling in the circle is 0.788
The estimated pi is 3.154



The fraction of samples falling in the circle is 0.788
The estimated pi is 3.152

Determine Value-at-Risk (VaR) and Expected Shortfall (ES)

Let r be the continuously compounded monthly return on a stock and W_0 be the initial wealth to be invested over the month. Assume that $r \sim N(0.01, 0.1^2)$ and $W_0 = \$10,000$.

Determine 5% VaR

Let W_1 be the end of month wealth and $L_1 \equiv W_1 - W_0$ be the profit of the investment, which can be negative. The monthly VaR on the W_0 investment with probability α is denoted by VaR_α such that

$$P(L_1 \leq VaR_\alpha) = \alpha.$$

Note that

$$L_1 = W_1 - W_0 = W_0 \cdot (e^r - 1).$$

Thus, we have

$$P\left(W_0 \cdot (e^r - 1) \leq VaR_\alpha\right) = \alpha.$$

It implies that

$$P(r \leq x) = \alpha,$$

where $x = \ln(\frac{VaR_\alpha}{W_0} + 1)$ or $VaR_\alpha = W_0 \cdot (e^x - 1)$. From the definition of quantile function (inverse CDF), we know $x = q_\alpha^r$. Thus, $VaR_\alpha = W_0 \cdot (e^{q_\alpha^r} - 1)$.

Therefore, to compute 5% VaR, we compute

1. $q_{5\%}^r$
2. $VaR_{5\%} = \$10,000 \cdot (e^{q_{5\%}^r} - 1)$

In [16]:

```
w0 = 10000
mu_r = 0.01
sigma_r = 0.1
alpha = 0.05

def VaR(alpha, w0 = w0, mu = mu_r, sigma = sigma_r):
    q_r = norm.ppf(alpha, mu, sigma)
    return w0 * (np.exp(q_r) - 1)

print("The 5% quantile of N({}, {:.2f}) is {:.3f}".format(mu_r, sigma_r**2, norm.
print("The 5% value at risk is {:.3f}".format(VaR(alpha)) )
```

The 5% quantile of N(0.01,0.01) is -3.280

The 5% value at risk is -1431.440

Determine 5% ES

Given the profit is not larger than VaR_α , the expected shortfall (ES) at α measures the conditional expected profit, i.e.,

$$ES_\alpha = E[L_1 | L_1 \leq VaR_\alpha] = \frac{\int_{-\infty}^{VaR_\alpha} x f_{L_1}(x) dx}{P(L_1 \leq VaR_\alpha)} = \alpha^{-1} \int_0^\alpha VaR_u du.$$

Next, we introduce several ways to compute ES_α , where $\alpha = 5\%$.

In [17]:

```
import scipy.integrate as integrate

integrate.quad(VaR, 0, alpha)
```

Out[17]: (-88.82667054220768, 8.17306187173017e-08)

In [18]:

```
ES = integrate.quad(VaR, 0, alpha)[0] / alpha
print("The 5% expected shortfall is {:.3f}".format (ES))
```

The 5% expected shortfall is -1776.533

Computations through MC simulations

The above conditional expectation can be approximated by its sample counterpart.

In [19]:

```
np.random.seed(123)
nsim = 5000
r = norm.rvs(loc = mu_r, scale = sigma_r, size = nsim)

X = w0 * (np.exp(r) - 1)
Y = (X <= VaR(alpha)) * 1

ES_sample = np.mean(Y * X)/np.mean(Y)
print("The simulated ES is %.3f" % ES_sample)
```

The simulated ES is -1759.723